

Hyperparameter Tuning for Deep Learning: From Zero to Hero

A Comprehensive, Beginner-Friendly Guide

Welcome! This notebook will teach you how to find the best hyperparameters for your deep learning models - a crucial skill that separates good models from great ones.

What You'll Learn

- What hyperparameters are and why they matter
- Manual tuning strategies and intuitions
- Automated search methods (Grid, Random, Bayesian)
- Modern tools (Optuna, Ray Tune)
- Best practices for efficient tuning

Table of Contents

Part	Topic	Key Concepts
I	Introduction	Parameters vs Hyperparameters
II	Key Hyperparameters	Learning rate, batch size, architecture
III	Manual Tuning	Intuition and heuristics
IV	Grid Search	Exhaustive but expensive
V	Random Search	Often better than grid!
VI	Bayesian Optimization	Smart, adaptive search
VII	Optuna Tutorial	Modern hyperparameter optimization
VIII	Best Practices	Efficient tuning strategies

```
# Install and import dependencies
!pip install -q numpy matplotlib torch scikit-learn optuna

import numpy as np
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torch.optim as optim
from sklearn.datasets import make_moons
from sklearn.model_selection import train_test_split
import warnings
warnings.filterwarnings('ignore')
```

```

np.random.seed(42)
torch.manual_seed(42)

plt.style.use('seaborn-v0_8-whitegrid')
plt.rcParams['figure.figsize'] = (12, 6)

print("Libraries loaded successfully!")

```

419.5/419.5 kB 6.9 MB/s eta 0
Libraries loaded successfully!

✓ Part I: Introduction to Hyperparameters

1.1 Parameters vs Hyperparameters

Parameters	Hyperparameters
Learned during training	Set before training
Weights, biases	Learning rate, batch size
Updated by optimizer	Chosen by you (or search)
Millions of them	Typically 10-50

Why Hyperparameters Matter

The same architecture with different hyperparameters can give:

- **90% accuracy** with good hyperparameters
- **60% accuracy** with bad hyperparameters
- **Complete failure** with terrible hyperparameters

Hyperparameter tuning is often more important than architecture design!

```

# Demonstrate impact of hyperparameters
X, y = make_moons(n_samples=500, noise=0.2, random_state=42)
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, r

X_train_t = torch.FloatTensor(X_train)
y_train_t = torch.LongTensor(y_train)
X_val_t = torch.FloatTensor(X_val)
y_val_t = torch.LongTensor(y_val)

class SimpleNet(nn.Module):
    def __init__(self, hidden_size=32):
        super().__init__()

```

```

        self.net = nn.Sequential(
            nn.Linear(2, hidden_size),
            nn.ReLU(),
            nn.Linear(hidden_size, hidden_size),
            nn.ReLU(),
            nn.Linear(hidden_size, 2)
        )

    def forward(self, x):
        return self.net(x)

def train_model(lr, hidden_size, epochs=100):
    torch.manual_seed(42)
    model = SimpleNet(hidden_size)
    optimizer = optim.Adam(model.parameters(), lr=lr)
    criterion = nn.CrossEntropyLoss()

    train_losses, val_accs = [], []
    for _ in range(epochs):
        model.train()
        optimizer.zero_grad()
        loss = criterion(model(X_train_t), y_train_t)
        loss.backward()
        optimizer.step()
        train_losses.append(loss.item())

        model.eval()
        with torch.no_grad():
            acc = (model(X_val_t).argmax(1) == y_val_t).float().mean().item()
            val_accs.append(acc)

    return train_losses, val_accs

# Test different learning rates
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

learning_rates = [0.0001, 0.001, 0.01, 0.1, 0.5]
colors = plt.cm.viridis(np.linspace(0, 1, len(learning_rates)))

for lr, color in zip(learning_rates, colors):
    losses, accs = train_model(lr, hidden_size=32)
    axes[0].plot(losses, color=color, label=f'LR={lr}')
    axes[1].plot(accs, color=color, label=f'LR={lr} (final: {accs[-1]:.1})

axes[0].set_xlabel('Epoch')
axes[0].set_ylabel('Training Loss')
axes[0].set_title('Impact of Learning Rate on Training', fontsize=12)
axes[0].legend()
axes[0].set_yscale('log')

axes[1].set_xlabel('Epoch')

```


```

axes[1].set_ylabel('Validation Accuracy')
axes[1].set_title('Same Architecture, Different LR = Very Different Resu
axes[1].legend()

plt.tight_layout()
plt.show()

print("\n" + "="*60)
print("KEY INSIGHT: Learning rate alone can make or break training!")
print("Finding good hyperparameters is crucial for success.")
print("="*60)

```



=====

KEY INSIGHT: Learning rate alone can make or break training!

Finding good hyperparameters is crucial for success.

=====

✓ Part II: Key Hyperparameters in Deep Learning

2.1 The Most Important Hyperparameters

Category	Hyperparameter	Typical Range	Impact
Optimization	Learning rate	1e-5 to 1e-1	Very High
	Batch size	16 to 512	High
	Optimizer	Adam, SGD, AdamW	Medium
	Weight decay	1e-5 to 1e-2	Medium
Architecture	Hidden layers	1 to 10+	High
	Hidden units	32 to 4096	Medium
	Dropout rate	0 to 0.5	Medium
	Activation	ReLU, GELU	Low-Medium

Category	Hyperparameter	Typical Range	Impact
Training	Epochs	10 to 1000	Medium
	LR schedule	Step, Cosine	Medium

```
# Visualize hyperparameter search space
fig, axes = plt.subplots(2, 2, figsize=(14, 10))

# Learning rate landscape
ax = axes[0, 0]
lrs = np.logspace(-5, 0, 50)
# Simulate typical accuracy vs LR curve
acc_curve = 0.5 + 0.4 * np.exp(-((np.log10(lrs) + 2.5)**2) / 1.5)
acc_curve += np.random.randn(50) * 0.02
ax.semilogx(lrs, acc_curve, 'b-', linewidth=2)
ax.fill_between(lrs, acc_curve - 0.05, acc_curve + 0.05, alpha=0.3)
ax.axvline(x=0.001, color='green', linestyle='--', label='Sweet spot')
ax.axvspan(0.0001, 0.01, alpha=0.2, color='green', label='Good range')
ax.set_xlabel('Learning Rate')
ax.set_ylabel('Validation Accuracy')
ax.set_title('Learning Rate vs Accuracy\n(There is usually a sweet spot!)')
ax.legend()

# Batch size tradeoffs
ax = axes[0, 1]
batch_sizes = [8, 16, 32, 64, 128, 256, 512]
train_time = [8, 5, 3, 2, 1.5, 1.2, 1] # Relative
generalization = [0.95, 0.94, 0.92, 0.90, 0.88, 0.85, 0.82] # Hypothetical

ax2 = ax.twinx()
ax.bar([str(b) for b in batch_sizes], train_time, alpha=0.7, color='blue')
ax2.plot([str(b) for b in batch_sizes], generalization, 'ro-', linewidth=2)
ax.set_xlabel('Batch Size')
ax.set_ylabel('Training Time (relative)', color='blue')
ax2.set_ylabel('Generalization', color='red')
ax.set_title('Batch Size Tradeoffs\n(Larger = faster but may generalize)')

# Model complexity vs performance
ax = axes[1, 0]
complexity = np.linspace(1, 100, 100)
train_acc = 0.5 + 0.5 * (1 - np.exp(-complexity/20))
val_acc = 0.5 + 0.45 * (1 - np.exp(-complexity/15)) - 0.003 * complexity
val_acc = np.maximum(val_acc, 0.3)

ax.plot(complexity, train_acc, 'b-', linewidth=2, label='Training Accuracy')
ax.plot(complexity, val_acc, 'r-', linewidth=2, label='Validation Accuracy')
ax.axvline(x=30, color='green', linestyle='--', label='Optimal complexity')
ax.fill_between(complexity[30:], train_acc[30:], val_acc[30:], alpha=0.3)
ax.set_xlabel('Model Complexity (hidden units)')
ax.set_ylabel('Accuracy')
ax.set_title('Bias-Variance Tradeoff\n(Too complex = overfitting)', font
```

```
ax.legend()
```

```
# Hyperparameter importance
```

```
ax = axes[1, 1]
```

```
params = ['Learning Rate', 'Batch Size', 'Depth', 'Width', 'Dropout', 'W
```

```
importance = [0.95, 0.7, 0.6, 0.5, 0.4, 0.35]
```

```
colors = plt.cm.RdYlGn(np.array(importance))
```

```
ax.barh(params, importance, color=colors)
```


```
ax.set_xlabel('Relative Importance')
```

```
ax.set_title('Hyperparameter Importance\n(Focus on what matters most!)',
```

```
ax.set_xlim(0, 1)
```

```
plt.tight_layout()
```

```
plt.show()
```



✓ Part III: Manual Tuning Strategies

3.1 The Learning Rate Finder

Start with learning rate - it's the most important!

Strategy: Gradually increase LR and watch the loss

- If loss **decreases** → LR is in good range
- If loss **explodes** → LR is too high
- If loss **barely moves** → LR is too low

```
# Learning Rate Finder implementation
def lr_finder(model_class, train_data, start_lr=1e-7, end_lr=1, num_step
    """Find optimal learning rate by gradually increasing it."""
    X_t, y_t = train_data
    torch.manual_seed(42)
    model = model_class()
    optimizer = optim.SGD(model.parameters(), lr=start_lr)
    criterion = nn.CrossEntropyLoss()

    # Exponential LR increase
    lr_mult = (end_lr / start_lr) ** (1 / num_steps)

    lrs, losses = [], []
    lr = start_lr
    best_loss = float('inf')

    for _ in range(num_steps):
        optimizer.zero_grad()
        loss = criterion(model(X_t), y_t)

        # Stop if loss explodes
        if loss.item() > 4 * best_loss or np.isnan(loss.item()):
            break

        best_loss = min(best_loss, loss.item())
        lrs.append(lr)
        losses.append(loss.item())

        loss.backward()
        optimizer.step()

        # Increase learning rate
        lr *= lr_mult
        for param_group in optimizer.param_groups:
            param_group['lr'] = lr

    return lrs, losses

# Run LR finder
lrs, losses = lr_finder(SimpleNet, (X_train_t, y_train_t))

fig, ax = plt.subplots(figsize=(12, 5))
```

```

ax.semilogx(lrs, losses, 'b-', linewidth=2)

# Find suggested LR (steepest descent point)
gradients = np.gradient(losses)
min_grad_idx = np.argmin(gradients[:len(gradients)//2]) # First half on
suggested_lr = lrs[min_grad_idx]


ax.axvline(x=suggested_lr, color='green', linestyle='--', linewidth=2,
          label=f'Suggested LR: {suggested_lr:.1e}')
ax.scatter([suggested_lr], [losses[min_grad_idx]], color='green', s=200,

ax.set_xlabel('Learning Rate')
ax.set_ylabel('Loss')
ax.set_title('Learning Rate Finder\n(Pick LR where loss decreases fastest)')
ax.legend(fontsize=12)
ax.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

print(f"\nSuggested learning rate: {suggested_lr:.1e}")
print("(Use a value slightly lower for stable training)")

```



Suggested learning rate: 2.7e-04
(Use a value slightly lower for stable training)

✓ Part IV: Grid Search

4.1 What is Grid Search?

Try **every combination** of hyperparameter values:

LR: [0.001, 0.01, 0.1]

Hidden: [32, 64, 128]

Grid:

(0.001, 32), (0.001, 64), (0.001, 128),

(0.01, 32), (0.01, 64), (0.01, 128),

(0.1, 32), (0.1, 64), (0.1, 128)

Total: $3 \times 3 = 9$ combinations

Pros and Cons

Pros	Cons
Complete coverage	Exponential growth
Reproducible	Wastes time on bad regions
Easy to parallelize	Requires good grid design

```
# Grid Search implementation
from itertools import product

def evaluate_hyperparams(lr, hidden_size, dropout=0.0, epochs=50):
    """Train and evaluate with given hyperparameters."""
    torch.manual_seed(42)

    model = nn.Sequential(
        nn.Linear(2, hidden_size),
        nn.ReLU(),
        nn.Dropout(dropout),
        nn.Linear(hidden_size, hidden_size),
        nn.ReLU(),
        nn.Dropout(dropout),
        nn.Linear(hidden_size, 2)
    )

    optimizer = optim.Adam(model.parameters(), lr=lr)
    criterion = nn.CrossEntropyLoss()

    for _ in range(epochs):
        optimizer.zero_grad()
        loss = criterion(model(X_train_t), y_train_t)
        loss.backward()
        optimizer.step()

    model.eval()
    with torch.no_grad():
```

```

val_acc = (model(X_val_t).argmax(1) == y_val_t).float().mean().i

return val_acc

# Define grid
lr_grid = [0.001, 0.01, 0.1]
hidden_grid = [16, 32, 64, 128]

# Run grid search
results = {}
print("Running Grid Search...")
for lr, hidden in product(lr_grid, hidden_grid):
    acc = evaluate_hyperparams(lr, hidden)
    results[(lr, hidden)] = acc
    print(f" LR={lr}, Hidden={hidden}: Accuracy={acc:.2%}")

# Find best
best_params = max(results, key=results.get)
print(f"\nBest: LR={best_params[0]}, Hidden={best_params[1]} → {results[

```

```

Running Grid Search...
LR=0.001, Hidden=16: Accuracy=85.00%
LR=0.001, Hidden=32: Accuracy=86.00%
LR=0.001, Hidden=64: Accuracy=86.00%
LR=0.001, Hidden=128: Accuracy=90.00%
LR=0.01, Hidden=16: Accuracy=88.00%
LR=0.01, Hidden=32: Accuracy=98.00%
LR=0.01, Hidden=64: Accuracy=99.00%
LR=0.01, Hidden=128: Accuracy=97.00%
LR=0.1, Hidden=16: Accuracy=96.00%
LR=0.1, Hidden=32: Accuracy=98.00%
LR=0.1, Hidden=64: Accuracy=99.00%
LR=0.1, Hidden=128: Accuracy=98.00%

```

```

Best: LR=0.01, Hidden=64 → 99.00%

```

```

# Visualize grid search results
fig, ax = plt.subplots(figsize=(10, 6))

# Create heatmap
acc_matrix = np.zeros((len(lr_grid), len(hidden_grid)))
for i, lr in enumerate(lr_grid):
    for j, hidden in enumerate(hidden_grid):
        acc_matrix[i, j] = results[(lr, hidden)]

im = ax.imshow(acc_matrix, cmap='RdYlGn', aspect='auto')
ax.set_xticks(range(len(hidden_grid)))
ax.set_yticks(range(len(lr_grid)))
ax.set_xticklabels(hidden_grid)
ax.set_yticklabels(lr_grid)
ax.set_xlabel('Hidden Size')
ax.set_ylabel('Learning Rate')

```


```

ax.set_title('Grid Search Results\n(Green = Better Accuracy)', fontsize=

# Add text annotations
for i in range(len(lr_grid)):
    for j in range(len(hidden_grid)):
        text = ax.text(j, i, f'{acc_matrix[i, j]:.1%}',
                        ha='center', va='center', fontsize=12, fontweight

plt.colorbar(im, label='Validation Accuracy')
plt.tight_layout()
plt.show()

```



✓ Part V: Random Search

5.1 Why Random Search?

Key Insight: For most problems, only a few hyperparameters really matter!

Random search explores more values of important hyperparameters by chance.

Research Finding (Bergstra & Bengio, 2012)

"Random search is more efficient than grid search in high-dimensional spaces"

With the same budget:

- Grid search: 9 points in a 3×3 grid
- Random search: 9 different values of each parameter!

```
# Visualize Grid vs Random search
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

np.random.seed(42)

# Grid search points
ax = axes[0]
grid_x = np.array([0.2, 0.5, 0.8])
grid_y = np.array([0.2, 0.5, 0.8])
grid_xx, grid_yy = np.meshgrid(grid_x, grid_y)
ax.scatter(grid_xx, grid_yy, s=200, c='blue', marker='s', label='Grid Se

# Show projections
for x in grid_x:
    ax.axvline(x=x, color='blue', alpha=0.3, linestyle='--')
for y in grid_y:
    ax.axhline(y=y, color='blue', alpha=0.3, linestyle='--')

ax.set_xlabel('Hyperparameter 1 (Important!)')
ax.set_ylabel('Hyperparameter 2 (Less Important)')
ax.set_title('Grid Search: Only 3 values of HP1 tested', fontsize=12)
ax.set_xlim(0, 1)
ax.set_ylim(0, 1)
ax.legend()


# Random search points
ax = axes[1]
rand_x = np.random.rand(9)
rand_y = np.random.rand(9)
ax.scatter(rand_x, rand_y, s=200, c='red', marker='o', label='Random Sea

# Show projections on x-axis (important parameter)
for x in rand_x:
    ax.axvline(x=x, color='red', alpha=0.3, linestyle='--')

ax.set_xlabel('Hyperparameter 1 (Important!)')
ax.set_ylabel('Hyperparameter 2 (Less Important)')
ax.set_title('Random Search: 9 different values of HP1 tested!', fontsiz
ax.set_xlim(0, 1)
ax.set_ylim(0, 1)
ax.legend()
```

```
plt.tight_layout()
plt.show()

print("\n" + "="*60)
print("KEY INSIGHT: Random search covers more unique values!")
print("If only HP1 matters, random search tests 9 values vs 3.")
print("="*60)
```



```
=====
KEY INSIGHT: Random search covers more unique values!
If only HP1 matters, random search tests 9 values vs 3.
=====
```

```
# Random Search implementation
def random_search(n_trials=20):
    """Perform random search over hyperparameter space."""
    results = []

    for trial in range(n_trials):
        # Sample hyperparameters randomly
        lr = 10 ** np.random.uniform(-4, -1) # Log-uniform: 0.0001 to 0.01
        hidden = int(2 ** np.random.uniform(4, 8)) # 16 to 256
        dropout = np.random.uniform(0, 0.5)

        acc = evaluate_hyperparams(lr, hidden, dropout)
        results.append({'lr': lr, 'hidden': hidden, 'dropout': dropout,

        if (trial + 1) % 5 == 0:
            print(f"Trial {trial+1}/{n_trials}: Best so far = {max(r['ac

    return results

# Run random search
print("Running Random Search...")
random_results = random_search(n_trials=20)
```

```
# Find best
best = max(random_results, key=lambda x: x['accuracy'])
print(f"\nBest found:")
print(f"  LR: {best['lr']:.4f}")
print(f"  Hidden: {best['hidden']}")
print(f"  Dropout: {best['dropout']:.2f}")
print(f"  Accuracy: {best['accuracy']:.2%}")
```

```
Running Random Search...
Trial 5/20: Best so far = 98.00%
Trial 10/20: Best so far = 99.00%
Trial 15/20: Best so far = 99.00%
Trial 20/20: Best so far = 99.00%
```

```
Best found:
  LR: 0.0703
  Hidden: 232
  Dropout: 0.40
  Accuracy: 99.00%
```

Part VI: Bayesian Optimization

6.1 The Smart Search

Bayesian optimization **learns from previous trials** to suggest better points!

How It Works

1. **Build a model** of the objective function (surrogate model)
2. **Use it to predict** promising regions
3. **Balance exploration vs exploitation**
4. **Update the model** after each trial

Key Concepts

- **Surrogate Model:** Gaussian Process or Tree Parzen Estimator
- **Acquisition Function:** Decides where to sample next
- **Expected Improvement:** Common acquisition function

✓ Part VII: Optuna Tutorial

7.1 Modern Hyperparameter Optimization

Optuna is a state-of-the-art hyperparameter optimization framework:

- Automatic pruning of bad trials
- Efficient sampling algorithms
- Easy to use API
- Great visualizations

```
import optuna
from optuna.visualization import plot_optimization_history, plot_param_i

def objective(trial):
    """Optuna objective function."""
    # Suggest hyperparameters
    lr = trial.suggest_float('lr', 1e-4, 1e-1, log=True)
    hidden_size = trial.suggest_int('hidden_size', 16, 256)
    dropout = trial.suggest_float('dropout', 0.0, 0.5)
    n_layers = trial.suggest_int('n_layers', 1, 4)

    # Build model
    torch.manual_seed(42)
    layers = []
    in_size = 2
    for _ in range(n_layers):
        layers.extend([
            nn.Linear(in_size, hidden_size),
            nn.ReLU(),
            nn.Dropout(dropout)
        ])
        in_size = hidden_size
    layers.append(nn.Linear(hidden_size, 2))
    model = nn.Sequential(*layers)

    # Train
    optimizer = optim.Adam(model.parameters(), lr=lr)
    criterion = nn.CrossEntropyLoss()

    for epoch in range(50):
        optimizer.zero_grad()
        loss = criterion(model(X_train_t), y_train_t)
        loss.backward()
        optimizer.step()

    # Pruning: Stop bad trials early
    model.eval()
    with torch.no_grad():
        val_acc = (model(X_val_t).argmax(1) == y_val_t).float().mean
    trial.report(val_acc, epoch)
```

```
        if trial.should_prune():
            raise optuna.TrialPruned()
        model.train()

    return val_acc

# Run Optuna optimization
study = optuna.create_study(direction='maximize',
                            pruner=optuna.pruners.MedianPruner())
study.optimize(objective, n_trials=30, show_progress_bar=True)

print("\n" + "="*60)
print("OPTUNA RESULTS")
print("="*60)
print(f"Best accuracy: {study.best_value:.2%}")
print(f"Best hyperparameters:")
for key, value in study.best_params.items():
    print(f"  {key}: {value}")
```



```
[I 2026-04-07 04:10:51,359] A new study created in memory with name: no-n
Best trial: 0. Best value: 0.99: 100% 30/30 [00:06<00:00, 7.91it/s]

[I 2026-04-07 04:10:52,353] Trial 0 finished with value: 0.99000000953674
[I 2026-04-07 04:10:52,783] Trial 1 finished with value: 0.97000002861022
[I 2026-04-07 04:10:53,354] Trial 2 finished with value: 0.99000000953674
[I 2026-04-07 04:10:53,618] Trial 3 finished with value: 0.87000000476837
[I 2026-04-07 04:10:53,774] Trial 4 finished with value: 0.93000000715255
[I 2026-04-07 04:10:53,790] Trial 5 pruned.
[I 2026-04-07 04:10:54,079] Trial 6 finished with value: 0.97000002861022
[I 2026-04-07 04:10:54,107] Trial 7 pruned.
[I 2026-04-07 04:10:54,535] Trial 8 finished with value: 0.98000001907348
[I 2026-04-07 04:10:54,552] Trial 9 pruned.
[I 2026-04-07 04:10:54,616] Trial 10 pruned.
[I 2026-04-07 04:10:54,657] Trial 11 pruned.
[I 2026-04-07 04:10:56,003] Trial 12 finished with value: 0.9599999785423
[I 2026-04-07 04:10:56,366] Trial 13 finished with value: 0.9800000190734
[I 2026-04-07 04:10:56,391] Trial 14 pruned.
[I 2026-04-07 04:10:56,416] Trial 15 pruned.
[I 2026-04-07 04:10:56,557] Trial 16 pruned.
[I 2026-04-07 04:10:57,281] Trial 17 finished with value: 0.9700000286102
[I 2026-04-07 04:10:57,310] Trial 18 pruned.
[I 2026-04-07 04:10:57,341] Trial 19 pruned.
[I 2026-04-07 04:10:57,366] Trial 20 pruned.
[I 2026-04-07 04:10:57,386] Trial 21 pruned.
[I 2026-04-07 04:10:57,410] Trial 22 pruned.
[I 2026-04-07 04:10:57,436] Trial 23 pruned.
[I 2026-04-07 04:10:58,135] Trial 24 finished with value: 0.9800000190734
[I 2026-04-07 04:10:58,160] Trial 25 pruned.
[I 2026-04-07 04:10:58,176] Trial 26 pruned.
[I 2026-04-07 04:10:58,199] Trial 27 pruned.
[I 2026-04-07 04:10:58,222] Trial 28 pruned.
[I 2026-04-07 04:10:58,248] Trial 29 pruned.
```

=====

OPTUNA RESULTS

=====

```
Best accuracy: 99.00%
Best hyperparameters:
  lr: 0.0009337124287063657
  hidden_size: 238
  dropout: 0.027291363063632668
  n_layers: 4
```

```
# Visualize Optuna results
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Optimization history
ax = axes[0]
trials = [t for t in study.trials if t.state == optuna.trial.TrialState.
values = [t.value for t in trials]
best_values = np.maximum.accumulate(values)
```


```

ax.plot(values, 'b.', alpha=0.5, label='Trial Accuracy')
ax.plot(best_values, 'r-', linewidth=2, label='Best So Far')
ax.set_xlabel('Trial')
ax.set_ylabel('Validation Accuracy')
ax.set_title('Optuna Optimization History', fontsize=12)
ax.legend()

# Parameter importance
ax = axes[1]
importances = optuna.importance.get_param_importances(study)
params = list(importances.keys())
values = list(importances.values())
colors = plt.cm.RdYlGn(np.array(values) / max(values))
ax.barh(params, values, color=colors)
ax.set_xlabel('Importance')
ax.set_title('Hyperparameter Importance\n(Based on optimization history)

plt.tight_layout()
plt.show()

```



Part VIII: Best Practices

8.1 Efficient Tuning Strategy

1. START SIMPLE
 - └ Use defaults, ensure code works
2. FIND LEARNING RATE
 - └ Use LR finder or coarse search

3. TUNE IMPORTANT PARAMS FIRST

└ LR, batch size, then architecture
4. USE RANDOM/BAYESIAN SEARCH

└ More efficient than grid search
5. USE EARLY STOPPING

└ Don't waste time on bad configs
6. SCALE UP GRADUALLY

└ Start small, increase budget for promising regions

8.2 Common Mistakes to Avoid

Mistake	Why It's Bad	Solution
Tuning on test set	Data leakage!	Use validation set
Too wide ranges	Wastes trials	Start narrow, expand
Ignoring interactions	Missing good combos	Use Bayesian optimization
Linear search for LR	Wrong scale	Use log scale
No reproducibility	Can't repeat results	Set random seeds

8.3 Summary Cheat Sheet

HYPERPARAMETER TUNING CHEAT SHEET
Learning Rate:
• Start: 1e-3 (Adam), 0.1 (SGD)
• Search: Log scale [1e-5, 1e-1]
• Use LR finder or warmup